FITZ WU (WU KE)

IANG holder • +852 8495 1754 (WhatsApp) • kewu2@ln.hk • github.com/WuKe29

Education

Lingnan University, Hong Kong **Sep 2025 – Nov 2026**
MSc Artificial Intelligence and Business Analytics — Distinction (CGPA 3.57/4.0)
Central South University of Forestry and Technology, Changsha **Sep 2020 – Jul 2024**
BEng Computer Science and Technology

Experience

SF Technology (SF Express Group) — AI Centre, Product Department **Jun 2026 – Sep 2026**
Product Manager Intern

- Owned the requirements definition and user guide for the AGV call-button unit within the AI Centre's logistics-robotics product line, turning functional requirements and on-site procedures into documentation used by frontline staff.
- Supported on-site deployment and testing of that unit: joined site surveys and on-site test runs, and fed findings back into the requirements; contributed to the cloud dashboard design documentation.
- Designed and developed a Chrome extension (Manifest V3, content scripts, DevTools) that exports structured data from cloud-based web pages, replacing a manual copy-and-paste step; Claude Code as the main AI tool.

Projects

Stock Analysis Agent — independent project *2026*

- Architected a multi-agent analysis pipeline (LangGraph ReAct + LangChain): deterministic code scans the universe, specialist agents each receive only their own read-only tool set (technical / news / filings), and a supervisor decides — nothing in the pipeline can trade. Shipped as a CLI, an MCP server (FastMCP) and a FastAPI + ECharts dashboard; long-term memory over SQLite FTS5 with a trigram tokenizer for Chinese.
- Unified the trust model — agents propose, code executes: 22 write-operation tools denied at build time, single-use confirmation tokens with expiry and replay protection, position / stop-loss limits enforced in code, and a TRADING_ENABLED gate on the order path.
- Caught a live failure in a repeated evaluation run: the chat agent called `submit_order` and placed a real sell order despite a prompt-level "confirm before ordering" rule. Confirmed it in the broker's order log, then removed all write tools from the chat agent and deleted the prompt rule.
- Built the evaluation stack: 64 pytest unit tests, 12 trajectory-level cases that assert which tools were actually called rather than what the reply claims, and an L3 replay that scores every past proposal against its forward N-day return; run-level observability tracks latency, tokens, tool usage and success rate.

Gait-Based Person Identification — course paper, team leader *2026*

- Led the study end to end — framing, evaluation protocol, modelling and interpretability analysis — coordinating four teammates' contributions and delivering the final report: Kuopio IMU dataset (51 subjects × 3 walking speeds, 47 after exclusions), 520 handcrafted features plus a 42-feature baseline, and a 1D-CNN on raw acceleration; both routes reach ≈99% closed-set accuracy.
- Ran the SHAP and Gini importance analysis: the pelvis was the most informative sensor and spectral features the strongest discriminator; dimensionality reduction to 48 features retained ≈99% accuracy, turning model output into a sensor-selection recommendation.

Skills

- **LLM & agents**: multi-agent orchestration with LangGraph (ReAct), LangChain, MCP client + FastMCP server, tool scoping and deny-lists, single-use confirmation tokens, LLM evaluation (unit, trajectory, decision replay) and observability, long-term memory (SQLite FTS5).
- **Programming & data**: Python (primary), JavaScript (Chrome Extension MV3), SQL (CTEs, window functions, joins); pandas, NumPy, SciPy, scikit-learn, PyTorch, SHAP; MongoDB (schema design).
- **Tools**: Git / GitHub, Docker, FastAPI, Jupyter, Linux; Claude Code, Cursor, Codex.

Additional

- SAS Certified Specialist: Machine Learning (2026); Blue Bridge Cup, Provincial Third Prize (2022).
- **Languages**: Mandarin (native); English (MSc taught in English); Cantonese (understands; speaking in progress).